

Non-Preemptive Flow with Rejections

x

Carnegie Mellon University

x

ICALP 2018

Job Scheduling Problem

- Jobs arrive online
 - Release time r_j , processing time p_j , weight w_j
- Single machine (this talk)
 - Finish time t_j : time job j finished by machine ($t_j \geq r_j + p_j$)
- Goal: minimize **total flow time**:

$$\sum_{\text{job } j} w_j \cdot (t_j - r_j) = \sum_{\text{job } j} w_j \cdot (\text{time } j \text{ stays in system})$$

- Online: do not know future jobs in advance
- Goal: be competitive with best **offline**
- Different policies: preemption, rejection

- Typical setting:
 - **Preemption**: (can switch jobs before finishing current job)
 - **No rejection**: must finish all jobs
- [BBCD 03] $n^{\Omega(1)}$ competitive lower bound (randomized)
- Simple lower bound (**no preemption**, deterministic)
 - $t = 0$: job A , $w_A = 1$, $p_A = 1$
 - As soon as algo processes A , release B with $w_B = M$, $p_B = 1/M$ (M is large)
 - ALG cannot preempt $A \implies ALG \geq M$
 - $OPT = O(1)$
- To be competitive (with offline), online needs different model

- [KP 00] Speed augmentation model:
 - ALG's machine is **faster by $(1 + \epsilon)$**
 \iff for ALG, $p_j \leftarrow p_j / (1 + \epsilon)$
 - Only need to be competitive with offline without speed augmentation
 - [CGKM 09] $\text{poly}(1/\epsilon)$ -competitive
 - Greedy algo
 - Proof: potential function
 - Also works for multiple, unrelated machines
 - [AGK 12] Proof via **dual fitting**
- [CDGK 18] Rejection model
 - **No preemption**: cannot switch jobs until current job done
 - **Rejection**: can reject ϵ **fraction total weight** of jobs
 - Competitive with offline **preemption, no rejection**
 - Intuition: rejecting “random” ϵ -fraction is like $(1 + \epsilon)$ -speed
- **Main result: $\text{poly}(1/\epsilon)$ -competitive in this model**
- Follow dual fitting framework of [AGK 12]

Dual Fitting Framework [AGK 12]

- Dual fitting: **systematic** way to study online algos
- For job j :
 - X := flow time of **HDF policy** for jobs 1 to $(j - 1)$
 - Y := flow time of **HDF policy** for jobs 1 to j
 - $\alpha_j := Y - X$
 - α_j is **marginal** increase in flow time from job j (for HDF)
 - Motivation: α_j represents a **dual variable** in LP relaxation
- $\sum_j \alpha_j$ is total flow time of HDF algo
- Thm (dual fitting): $OPT \geq \sum_j \alpha_j - ALG$
- Goal: $ALG \leq (1 - \epsilon^3) \cdot \sum_j \alpha_j$
 $\implies 1/\epsilon^3$ -competitive
- Intuition:
 - 1 Make ALG “almost-HDF”: $ALG \leq (1 + \epsilon^3) \cdot \sum_{\text{accepted } j} \alpha_j$
 - 2 Reject jobs with large α_j (can compute α_j online!)
 $\implies \sum_{\text{accepted } j} \alpha_j \leq (1 - 2\epsilon^3) \cdot \sum_{\text{all } j} \alpha_j$

(1) Make ALG “almost-HDF”

- Problem: online model is **no preemption**
- Fix: whenever preempt, reject instead
- Goal: maintain almost-HDF, but preempt only ϵ -fraction
- Algorithm outline:
 - For current job j , time t : consider **lighter** jobs that arrive **later** than j
 - $r_j \leq r_{j'} \leq t$ and $w_{j'} < w_j$
 - HDF should have preempted j for some j'
 - If total weight of such jobs exceeds w_j/ϵ , then preempt j
 - **Charge** preemption of j to the arrivals of these jobs
 \implies preempt $\leq \epsilon$ -fraction of total weight
- Thm: if always process a current job j such that

$$\sum_{\substack{\text{unfinished} \\ \text{lighter jobs } j'}} w_{j'} \leq w_j/\epsilon,$$

i.e., “almost-HDF”, then $ALG \leq \sum_{\text{acc } j} \alpha_j + \sum_j w_j p_j / \epsilon$

(2) Reject Jobs with Large α_j

- Goal: reject jobs such that $\sum_{\text{rej } j} \alpha_j \geq \text{poly}(\epsilon) \cdot \sum_j \alpha_j$ and $\sum_{\text{rej } j} w_j \leq \epsilon \cdot \sum_j w_j$
- Recall: can compute α_j values online
- Randomized achieves it in expectation
- Deterministic is harder: most technically involved part
 - Bucket into geometric increasing classes based on α_j, w_j, p_j
 - Charge to class with largest geometric power

Open Questions

- Total flow time $\sum_{\text{job } j} w_j \cdot (t_j - r_j)$ is “ ℓ_1 -norm”
- Can we prove results for ℓ_p norm (constant p)?
- [AGK 12] Speed augmentation setting: same dual fitting technique solves ℓ_p norm
- What about ℓ_∞ norm? **name for it?**